

# Learning Breakout using NEAT

James T. Murphy III<sup>1</sup>   Skyler Thomas<sup>2</sup>

The University of Texas at Austin

December 12, 2017

---

<sup>1</sup>jamesmurphy@math.utexas.edu

<sup>2</sup>skylerthomas0@gmail.com

# A neat application of a neat algorithm

In 2015, YouTuber SethBling trained a neural network to successfully beat the first level of Super Mario World [2].



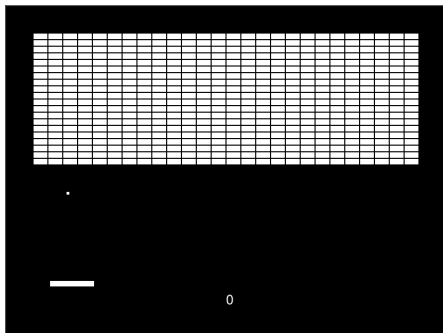
This was done using the Neural Evolution of Augmenting Technologies algorithm (NEAT) from the 2002 paper by Stanley and Mikkulainen [1].

# NEAT, the algorithm

NEAT is a type of genetic algorithm that follows the biological metaheuristics of “fitness begets evolution”.

1. Model neural network nodes and weights as genes.
2. Randomly mutate genomes (add/delete node, add/delete connection, change connection weight, etc.).
3. Rank individuals by fitness.
4. Allow best performing to mate using genetic crossover algorithm.
5. Repeat until fitness threshold is reached.

# Breakout (AKA BrickBreaker)



We use the NEAT algorithm to train a network to play Breakout. The fitness of a genome is it's final score (out of 520) after playing a game of Breakout.

# Results

Qualitatively, it appears there are three eras:

1. Pre-ball-tracking era: The network does nothing or makes small and arbitrary movements.
2. Ball-tracking era: It is apparent the paddle is tracking the ball, but it may make mistakes or infinite loop without clearing some remaining blocks.
3. Aiming era: The network actively moves to and aims the ball to clear all 520 blocks.


# Results

Qualitatively, it appears there are three eras:

1. Pre-ball-tracking era: The network does nothing or makes small and arbitrary movements.
2. Ball-tracking era: It is apparent the paddle is tracking the ball, but it may make mistakes or infinite loop without clearing some remaining blocks.
3. Aiming era: The network actively moves to and aims the ball to clear all 520 blocks.

Demo time!


# Results



**Figure:** *A successful training experiment.* Generations of 200 individuals took an average time of 2.835 sec to evaluate on an Intel Core i7-4790K @ 4.00GHz in parallel on 6 cores. Total runtime  $\approx$  17.5 minutes.

# The winning network

The most fit network is extremely simple. It has no hidden nodes and only a few connections.



This network is far too small to have memorized a sequence of inputs, it's actually tracking the ball.



# Discussion

- ▶ Full board information too many inputs for success.
- ▶ Needed to give history of ball location for success.
- ▶ Needed to train with deterministic initial condition.
- ▶ Training prone to stagnation in a local maximum fitness.
- ▶ Success is very sensitive to tiny changes in hyperparameters.
- ▶ Originally attempted to do Tetris, which was too complicated for NEAT with the current setup.

# Future

- ▶ Change setup:
  1. Randomly generate initial board.
  2. Make score based on 5 seconds of gameplay instead of full game.
  3. Optionally, make score weighted by the amount of time it takes to achieve it to encourage clearing blocks *quickly*.
- ▶ Compare NEAT vs. other evolutionary algorithms vs. non-evolutionary algorithms like Deep Q-learning.



[1] K. O. Stanley and R. Miikkulainen.

Evolving neural networks through augmenting topologies.

*Evolutionary computation*, 10(2):99–127, 2002.



[2] SethBling.

Mari/o machine learning for video games.

<https://www.youtube.com/watch?v=qv6UV0Q0F44>, June 2015.



[3] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves,

I. Antonoglou, D. Wierstra, and M. Riedmiller.

Playing atari with deep reinforcement learning.

*arXiv preprint arXiv:1312.5602*, 2013.



[4] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness,

M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland,

G. Ostrovski, et al.

Human-level control through deep reinforcement learning.

*Nature*, 518(7540):529–533, 2015.